

Hashing

Prof. Ki-Hoon Lee
Dept. of Computer Engineering
Kwangwoon University

Dictionaries

- Collection of pairs.
 - (key, element)
 - Pairs have different keys.
- Operations.
 - Get(theKey)
 - Delete(theKey)
 - Insert(theKey, theElement)

Hash Tables

- Worst-case time for **Get**, **Insert**, and **Delete** is **$O(n)$** .
 - **n**: the number of pairs in a dictionary
- Expected time is **$O(1)$** .

Ideal Hashing

- Uses a 1D array (or table) $\text{table}[0:b-1]$.
 - Each position of this array is a bucket.
 - b is the number of buckets in the table.
 - A bucket can normally hold only one dictionary pair.
- Uses a hash function f that converts each key k into an index in the range $[0, b-1]$.
 - $f(k)$ is the home bucket for key k .
- Every dictionary pair $(\text{key}, \text{element})$ is stored in its home bucket $\text{table}[f(\text{key})]$.

Ideal Hashing Example

- Pairs are: (22,a), (33,c), (3,d), (73,e), (85,f).
- Hash table is $\text{table}[0:7]$, $b = 8$.
- Hash function is $\text{key}/11$.
- Pairs are stored in table as below:

(3,d)		(22,a)	(33,c)			(73,e)	(85,f)
[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]

- **Get**, **Insert**, and **Delete** take $O(1)$ time.

What Can Go Wrong?

(3,d)		(22,a)	(33,c)			(73,e)	(85,f)
[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]

- Where does (26,g) go?
- Keys that have the same home bucket are **synonyms**.
 - 22 and 26 are synonyms with respect to the hash function that is in use.
- The home bucket for (26,g) is already occupied.

What Can Go Wrong?

(3,d)		(22,a)	(33,c)			(73,e)	(85,f)
-------	--	--------	--------	--	--	--------	--------

- A **collision** occurs when the home bucket for a new pair is occupied by a pair with a different key.
- An **overflow** occurs when there is no space in the home bucket for the new pair.
- When a bucket can hold only one pair, collisions and overflows occur together.
- Need a method to handle overflows.

Hash Table Issues

- Choice of hash function.
- Overflow handling method.
- Size (number of buckets) of hash table.

Hash Functions

- Two parts:
 - Convert key into a nonnegative integer in case the key is not an integer.
 - Done by the function `hash()`.
 - Map an integer into a home bucket.
 - $f(k)$ is an integer in the range $[0, b-1]$, where b is the number of buckets in the table.

String to Integer

- Each character is **1** byte long.
- An **int** is **4** bytes.
- A **2** character string **s** may be converted into a unique **4** byte non-negative **int** using the code:

```
int answer = s.at(0);
```

```
answer = (answer << 8) + s.at(1);
```

- Strings that are longer than **3** characters do not have a unique **non-negative int** representation.

String to Nonnegative Integer

```
template<>
class hash<string>
{
    public:
    size_t operator()(const string theKey) const
    { // Convert theKey to a nonnegative integer.
        unsigned long hashValue = 0;
        int length = (int) theKey.length();
        for (int i = 0; i < length; i++)
            hashValue = 5 * hashValue +
                        theKey.at(i);

        return size_t(hashValue);
    }
};
```

Map into a Home Bucket

(3,d)		(22,a)	(33,c)			(73,e)	(85,f)
[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]

- Most common method is by division.
 $\text{homeBucket} = \text{hash}(\text{theKey}) \% \text{divisor};$
- divisor equals the number of buckets b .
- $0 \leq \text{homeBucket} < \text{divisor} = b$

Uniform Hash Function

(3,d)		(22,a)	(33,c)			(73,e)	(85,f)
[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]

- Let **keySpace** be the set of all possible keys.
- A **uniform hash function** maps the keys in **keySpace** into buckets such that approximately the same number of keys get mapped into each bucket.

Uniform Hash Function

(3,d)		(22,a)	(33,c)			(73,e)	(85,f)
[0]	[1]	[2]	[3]	[4]	[5]	[6]	[7]

- Equivalently, the probability that a randomly selected key has bucket i as its home bucket is $1/b$, $0 \leq i < b$.
- A uniform hash function minimizes the likelihood of an overflow when keys are selected at random.

Hashing by Division

- $\text{keySpace} = \text{all ints}$.
- For every b , the number of ints that get mapped (hashed) into bucket i is approximately $2^{32}/b$.
- Therefore, the division method results in a uniform hash function when $\text{keySpace} = \text{all ints}$.
- In practice, keys tend to be correlated.
- So, the choice of the divisor b affects the distribution of home buckets.

Selecting the Divisor

- Because of this correlation, applications tend to have a bias towards keys that map into odd integers (or into even ones).
- When the divisor is an even number, odd integers hash into odd home buckets and even integers into even home buckets.
 - $20\%14 = 6$, $30\%14 = 2$, $8\%14 = 8$
 - $15\%14 = 1$, $3\%14 = 3$, $23\%14 = 9$
- The bias in the keys results in a bias toward either the odd or even home buckets.

Selecting the Divisor (cont.)

- When the divisor is an odd number, odd (even) integers may hash into any home.
 - $20\%15 = 5$, $30\%15 = 0$, $8\%15 = 8$
 - $15\%15 = 0$, $3\%15 = 3$, $23\%15 = 8$
- The bias in the keys does not result in a bias toward either the odd or even home buckets.
- Better chance of uniformly distributed home buckets.
- So do not use an even divisor.

Selecting the Divisor (cont.)

- Similar biased distribution of home buckets is seen, in practice, when the divisor is a multiple of prime numbers such as 3, 5, 7, ...
- The effect of each prime divisor p of b decreases as p gets larger.
- Ideally, choose b so that it is a prime number.
- Alternatively, choose b so that it has no prime factor smaller than 20.

Overflow Handling

- An overflow occurs when the home bucket for a new pair (key, element) is full.
- We may handle overflows by:
 - Search the hash table in some systematic fashion for a bucket that is not full.
 - Linear probing (linear open addressing).
 - Quadratic probing.
 - Random probing.
 - Eliminate overflows by permitting each bucket to keep a list of all pairs for which it is the home bucket.
 - Array linear list.
 - Chain.

Linear Probing

- When inserting a new pair whose key is k , we search the hash table buckets in the order,
 $ht[(h(k) + i) \% b]$, $0 \leq i \leq b-1$
 - ht : hash table
 - h : hash function
 - b : the number of buckets
- This search terminates when we reach the first unfilled bucket and the new pair is inserted into this bucket
- In case no such bucket is found, the hash table is full and it is necessary to increase the table size

Linear Probing



Put(N,7)

Value	Key
RID	R
0	B
1	O
2	E
3	P
4	V
5	L
6	X
7	N
8	K
9	M

Hash(key)

0
2
3
3
9
10
0
1
9
0

First
Empty
Slot ?

	KEY	VAL
0	B	0
1	X	6
2	O	1
3	E	2
4	P	3
5	N	7
6		
7		
8		
9	V	4
10	L	5

%M

Linear Probing Example

insert(76) insert(93) insert(40) insert(47) insert(10) insert(55)
 $76\%7 = 6$ $93\%7 = 2$ $40\%7 = 5$ $47\%7 = 5$ $10\%7 = 3$ $55\%7 = 6$

0	
1	
2	
3	
4	
5	
6	76

probes: 1

0	
1	
2	93
3	
4	
5	
6	76

1

0	
1	
2	93
3	
4	
5	40
6	76

1

0	47
1	
2	93
3	
4	
5	40
6	76

3

0	47
1	
2	93
3	10
4	
5	40
6	76

1

0	47
1	55
2	93
3	10
4	
5	40
6	76

3

Search with Linear Probing

◆ Consider a hash table A that uses linear probing

◆ **get(k)**

- We start at cell $h(k)$
- We probe consecutive locations until one of the following occurs
 - ◆ An item with key k is found, or
 - ◆ An empty cell is found, or
 - ◆ N cells have been unsuccessfully probed

Algorithm *get(k)*

$i \leftarrow h(k)$

$p \leftarrow 0$

repeat

$c \leftarrow A[i]$

if $c = \emptyset$

return *null*

else if $c.key() = k$

return $c.element()$

else

$i \leftarrow (i + 1) \bmod N$

$p \leftarrow p + 1$

until $p = N$

return *null*

Program 8.4: Linear probing

```
=====
template <class K, class E>
pair<K,E>* LinearProbing<K,E>::Get(const K& k)
{
    // Search the linear probing hash table ht
    // (each bucket has exactly one slot) for k.
    // If a pair with this key is found,
    // return a pointer to this pair;
    // otherwise, return NULL.
    int i = h(k);    // home bucket
    int j;
    for (j = i; ht[j] && ht[j]->first != k;) {
        j = (j + 1) % b;    // treat the table as circular
        if (j == i) return NULL; // back to start point
    }
    if (ht[j]->first == k) return ht[j];
    return NULL;
}
=====
```


Performance of Linear Probing

0	4				8				12				16			
34	0	45				6	23	7			28	12	29	11	30	33

- Worst-case find/insert/erase time is $O(n)$, where n is the number of pairs in the table.
- This happens when all pairs are in the same cluster.

Expected Performance

0	4				8				12				16				
34	0	45				6	23	7				28	12	29	11	30	33

- α = loading density = (the number of pairs)/b.
 - $\alpha = 12/17$.
- S_n = expected number of buckets examined in a successful search when n is large
- U_n = expected number of buckets examined in an unsuccessful search when n is large
- Time to insert governed by U_n .

Expected Performance

- $S_n \approx \frac{1}{2}(1 + 1/(1 - \alpha))$
- $U_n \approx \frac{1}{2}(1 + 1/(1 - \alpha)^2)$
- Note that $0 \leq \alpha \leq 1$.

<i>alpha</i>	S_n	U_n
<i>0.50</i>	1.5	2.5
<i>0.75</i>	2.5	8.5
<i>0.90</i>	5.5	50.5

$\alpha \leq 0.75$ is
recommended.

Hash Table Design

- Performance requirements are given, determine maximum permissible loading density.
- We want a successful search to make no more than 10 compares (expected).
 - $S_n \approx \frac{1}{2}(1 + 1/(1 - \alpha))$
 - $\alpha \leq 18/19$
- We want an unsuccessful search to make no more than 13 compares (expected).
 - $U_n \approx \frac{1}{2}(1 + 1/(1 - \alpha)^2)$
 - $\alpha \leq 4/5$
- So $\alpha \leq \min\{18/19, 4/5\} = 4/5$.

Hash Table Design


- Dynamic resizing of table.
 - Whenever loading density exceeds threshold ($4/5$ in our example), rehash into a table of approximately twice the current size.
- Fixed table size.
 - $\alpha = (\text{the number of pairs})/b$
 - Know maximum number of pairs.
 - No more than 1000 pairs.
 - Loading density $\leq 4/5 \Rightarrow b \geq 5/4 * 1000 = 1250$.
 - Pick b (equal to **divisor**) to be a prime number or an odd number with no prime divisors smaller than 20 .

Linear List of Synonyms

- Each bucket keeps a linear list of all pairs for which it is the home bucket.
- The linear list may or may not be sorted by key.
- The linear list may be an array linear list or a chain (i.e., linked list).

Sorted Chains

- Put in pairs whose keys are
6, 12, 34, 29,
28, 11, 23, 7, 0,
33, 30, 45
- Home bucket =
key % 17.



Program 8.5:Chain search

=====

```
template <class K, class E>
pair<K,E>* Chaining<K,E>::Get(const K& k)
{
    // Search the chained hash table ht for k.
    // If a pair with this key is found,
    // return a pointer to this pair;
    // otherwise, return NULL.
    int i = h(k);    // home bucket
    // search the chain ht[i]
    for (ChainNode<pair<K,E> >* current = ht[i];
         current;
         current = current->link)
        if (current->data.first == k)
            return &current->data;
    return NULL;
}
```

=====

STL unordered_map (hash_map)

```
#include <iostream>
#include <string>
#include <unordered_map>

int main()
{
    // Create an unordered_map of three strings (that map to strings)
    std::unordered_map<std::string, std::string> u = {
        {"RED", "#FF0000"},
        {"GREEN", "#00FF00"},
        {"BLUE", "#0000FF"}
    };

    // Iterate and print keys and values of unordered_map
    for( const auto& n : u ) {
        std::cout << "Key:[" << n.first << "]" Value:[" << n.second << "]\n";
    }

    // Add two new entries to the unordered_map
    u["BLACK"] = "#000000";
    u["WHITE"] = "#FFFFFF";

    // Output values by key
    std::cout << "The HEX of color RED is:[" << u["RED"] << "]\n";
    std::cout << "The HEX of color BLACK is:[" << u["BLACK"] << "]\n";

    return 0;
}
```

Output:

```
Key:[RED] Value:[#FF0000]  
Key:[BLUE] Value:[#0000FF]  
Key:[GREEN] Value:[#00FF00]  
The HEX of color RED is:[#FF0000]  
The HEX of color BLACK is:[#000000]
```

auto specifier (since C++11)

- For variables, specifies that the type of the variable that is being declared will be automatically deduced from its initializer.

```
auto d = 5.0; // 5.0 is a double literal, so d will be type double
auto i = 1 + 2; // 1 + 2 evaluates to an integer, so i will be type int
```

```
int add(int x, int y)
{
    return x + y;
}
```

```
int main()
{
    auto sum = add(5, 6); // add() returns an int, so sum will be type int
    return 0;
}
```

auto (자동범위변수)

- 일반적인 지역 변수 형태로 블록 안에서만 유효하며 블록의 실행이 끝나면 소멸
- 스택에 메모리 할당
- 일반적으로 C에서 auto 키워드는 생략되어있음. 즉 아무 표시하지 않은 변수는 auto와 같은 의미.
- C++에서 auto 키워드를 사용할 경우 "자동 타입 추론"이라는 완전히 다른 의미를 가지게 되므로 주의할것.